

Bootstrap methods Part 2

Bootstrap:

Example

Can a moth remember what it learned as a caterpillar? (Blackiston et al., 2008, PLoS ONE)

We consider the odds ratio, where the odds (of success) in a group is the fraction of the proportion of success divided by the proportion of failure.

Air choice / Group	Treatment	Control
<i>Clear air</i>	32	25
<i>Specific odor</i>	9	21
Total	41	46

```
library(bootstrap)
B = 1000

treatment = c(rep(1,32), rep(0,9))
control = c(rep(1,25), rep(0,21))
set.seed(1)

pt = bootstrap(x=treatment, nboot=B, theta=mean)$thetastar
pc = bootstrap(x=control, nboot=B, theta=mean)$thetastar

OR = (pt/(1-pt))/(pc/(1-pc))
sd(OR)
```


```
[1] 2.01208
```

```
quantile(OR, c(0.025,0.975))
```

```
      2.5%    97.5%  
1.193182 8.913580
```

```
hist(OR, probability=T, breaks=20, col="lightyellow")  
abline(v=sd(OR), col="violet", lwd=2)
```

Histogram of OR



```
# Define the statistic  
proportion = function(x) return(sum(x==1)/length(x))  
set.seed(1)
```

```
pt = bootstrap(x=treatment, nboot=B, theta=proportion)$thetastar
pc = bootstrap(x=control, nboot=B, theta=proportion)$thetastar
OR = (pt/(1-pt))/(pc/(1-pc))
sd(OR)
```

```
[1] 2.01208
```

Regression analysis

The Duncan data frame has 45 rows and 4 columns. Data on the prestige and other characteristics of 45 U.S. occupations in 1950.

```
library(car)
library(MASS)
library(simpleboot)
data(Duncan)
attach(Duncan)

head(Duncan)
```

	type	income	education	prestige
accountant	prof	62	86	82
pilot	prof	72	76	83
architect	prof	75	92	90
author	prof	55	90	76
chemist	prof	64	86	90
minister	prof	21	84	87


```
model.dun1 = lm(prestige ~ education, data=Duncan)
coef(summary(model.dun1))
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.2839995	5.09306471	0.05576201	9.557897e-01
education	0.9019958	0.08455492	10.66757246	1.170879e-13

```
summary(model.dun1)$r.squared
```

```
[1] 0.7257602
```

```
plot(Duncan$education, Duncan$prestige)
abline(model.dun1)
```



```
model.dun2 = lm(prestige ~ income + education, data=Duncan)
coef(summary(model.dun2))
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.0646629	4.27194117	-1.419650	1.630896e-01
income	0.5987328	0.11966735	5.003310	1.053184e-05
education	0.5458339	0.09825264	5.555412	1.727192e-06

```
summary(model.dun2)$fstatistic
```

value	numdf	dendf
101.2162	2.0000	42.0000

```
summary(model.dun2)$adj.r.squared
```

```
[1] 0.8199912
```

Bootstrapping with package simpleboot

```
library(simpleboot)
summary(lm.boot(lm(prestige ~ income + education), R=B, rows=T))
```

```
BOOTSTRAP OF LINEAR MODEL (method = rows)
```

```
Original Model Fit
```

```
-----
```

```
Call:
```

```
lm(formula = prestige ~ income + education)
```

```
Coefficients:
```

(Intercept)	income	education
-6.0647	0.5987	0.5458

```
Bootstrap SD's:
```

(Intercept)	income	education
3.0776095	0.1682817	0.1378040

A **robust** regression model:

```
model.dun3 = rlm(prestige ~ income + education, data=Duncan)
coef(summary(model.dun3))
```

	Value	Std. Error	t value
(Intercept)	-7.1107028	3.88131509	-1.832034
income	0.7014493	0.10872497	6.451593
education	0.4854390	0.08926842	5.437970

```
B = 1000
```

```
rrpair = function(index, xdata){
  rlm(prestige ~ income + education, data=xdata[index,])$coefficients
}
set.seed(1)
```

```
rbet = bootstrap(x=1:45, B, theta=rrpair, xdata=Duncan)$thetastar  
sd(rbet[1,])
```

```
[1] 3.046292
```

```
sd(rbet[2,])
```

```
[1] 0.1787839
```

```
sd(rbet[3,])
```

```
[1] 0.1392966
```

Bootstrapping residuals

```
library(bootstrap)  
data(Duncan)  
  
regres = function(x, beta, xdata){  
  sprestige = beta[1] + beta[2]*xdata$income + beta[3]*xdata$education + x  
  
  return(coef(lm(sprestige ~ income + education, data=xdata)))  
}  
  
model = lm(prestige ~ income + education, data=Duncan)  
res = model$residuals  
beta.h = coef(model)  
  
B = 1000  
set.seed(1)  
bet.res = bootstrap(x=res, B, regres, beta=beta.h, xdata=Duncan)$thetastar  
  
sd(bet.res[1,])
```

```
[1] 4.100723
```

```
sd(bet.res[2,])
```

```
[1] 0.1141664
```

```
sd(bet.res[3,])
```

```
[1] 0.0979109
```

Bootstrapping residuals (with package simpleboot)

```
library(simpleboot)

summary(lm.boot(lm(prestige ~ income + education), R=B, rows=F))
```

```
BOOTSTRAP OF LINEAR MODEL (method = residuals)
```

```
Original Model Fit
```

```
-----
```

```
Call:
```

```
lm(formula = prestige ~ income + education)
```

```
Coefficients:
```

(Intercept)	income	education
-6.0647	0.5987	0.5458

```
Bootstrap SD's:
```

(Intercept)	income	education
4.10436442	0.11725724	0.09643463

Bootstrap regression with library car

```
N = 50
sd = 0.5
x = rnorm(N)
y = 10 * x + sd * rnorm(N)^2
datos = data.frame(y, x)
```

```
library(car)

modeloB = lm(y ~ x, datos)
betahat.boot = Boot(modeloB, R=2000)
summary(betahat.boot) # default summary
```

```
Number of bootstrap replications R = 2000
```


	original	bootBias	bootSE	bootMed
(Intercept)	0.41846	-0.0036317	0.072978	0.41159
x	9.96203	0.0066940	0.093776	9.97080

```
confint(betahat.boot)
```

Bootstrap bca confidence intervals

	2.5 %	97.5 %
(Intercept)	0.3055779	0.6225411
x	9.7187426	10.1081584

```
hist(betahat.boot)
```

With residuals:

```
betahat.boot2 = Boot(modeloB, method = 'residual', R=2000)
summary(betahat.boot2) # default summary
```


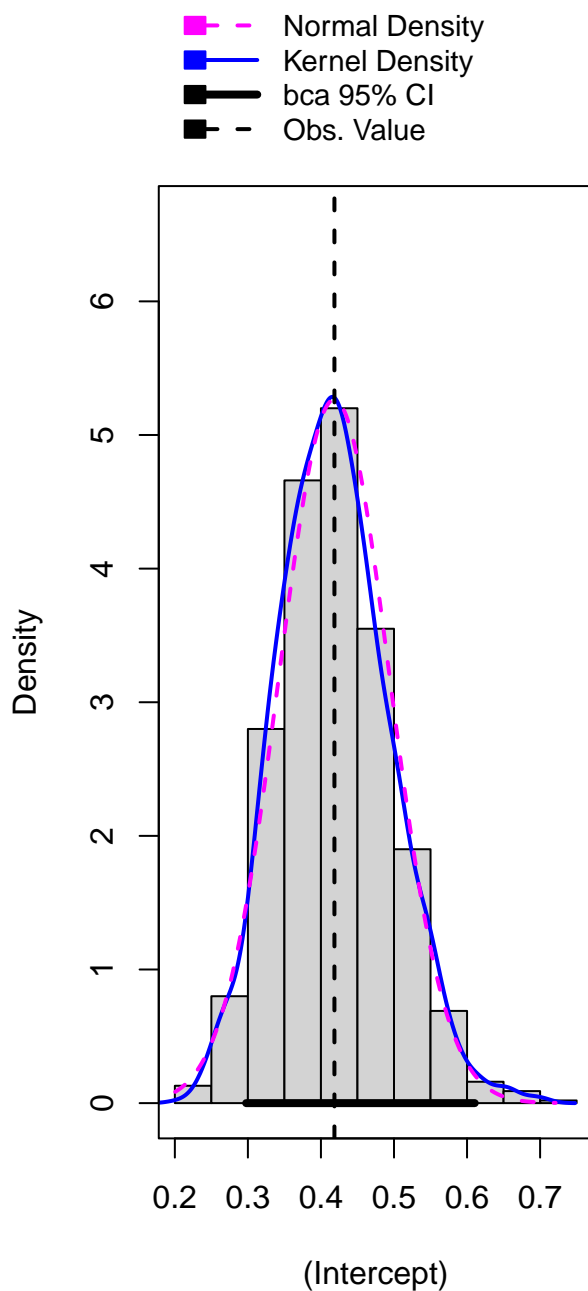
```
Number of bootstrap replications R = 2000
      original    bootBias  bootSE bootMed
(Intercept)  0.41846 -0.00023121 0.075772 0.41544
x            9.96203  0.00271319 0.080113 9.96321
```

```
confint(betahat.boot2)
```

Bootstrap bca confidence intervals

	2.5 %	97.5 %
(Intercept)	0.2976693	0.6105535
x	9.8018166	10.1241178

```
hist(betahat.boot2)
```



Time series

```
beta.real = 0.5

set.seed(123)
x = c(rnorm(1))

for(i in 2:150) x = c(x, x[i-1]*beta.real + rnorm(1))

(model = arima(x, order=c(1,0,0), include.mean=F))
```

```
Call:
arima(x = x, order = c(1, 0, 0), include.mean = F)


Coefficients:
      ar1
    0.4816
s.e.  0.0711

sigma^2 estimated as 0.8958:  log likelihood = -204.72,  aic = 413.44
```


```
par(mfrow=c(1,2))

plot(x, type="l", main="Simulated ts")
acf(x)
```

Simulated ts



Series x



```
library(bootstrap)

ar1fit = function(x,beta){
  y = arima.sim(n=150, list(ar=beta), innov=x)
  return(coef(arima(y, order=c(1,0,0), include.mean=F)))
}


res = model$residuals
beta.h = coef(model)

set.seed(1)
```

```
bet.res = bootstrap(x=res, 1000, ar1fit, beta=beta.h)$thetastar  
sd(bet.res)
```

```
[1] 0.07260983
```

```
hist(bet.res, probability=T)  
abline(v=beta.h, col="violet")
```



```
library(boot)
```

```
ar1 = function(ts) coef(arima(ts, c(1,0,0), include.mean=F))

(beta.bl = tsboot(x, statistic=ar1, R=1000, sim="fixed", l=5))
```

BLOCK BOOTSTRAP FOR TIME SERIES

Fixed Block Length of 5

Call:


```
tsboot(tseries = x, statistic = ar1, R = 1000, l = 5, sim = "fixed")
```

Bootstrap Statistics :

	original	bias	std. error
t1*	0.4815887	-0.09785202	0.07804321

```
hist(beta.bl$t, probability=T)
abline(v=beta.h, col="violet")
```

Histogram of beta.bl\$t



Confidence intervals

By assuming an exponential distribution, the exact confidence interval for the parameter (or for the mean) can be calculated. See:

en.wikipedia.org/wiki/Exponential_distribution

$$IC_{\frac{1}{\lambda}} = \left[\frac{\bar{x} \cdot 2n}{\chi_{2n, 1 - \frac{\alpha}{2}}^2}; \frac{\bar{x} \cdot 2n}{\chi_{2n, \frac{\alpha}{2}}^2} \right]$$

```
library(bootstrap)

data(aircondit, package="boot")
m = mean(aircondit$hours)
s = sd(aircondit$hours)
n = length(aircondit$hours)

# Exact CI (exponential)
c(m*2*n/qchisq(0.975, df=2*n), m*2*n/qchisq(0.025, df=2*n))
```

```
[1] 65.89765 209.17415
```

```
# Asymptotic CI
m + s/sqrt(n)*qnorm(0.975)*c(-1,1)
```

```
[1] 31.00421 185.16246
```

```
m + s/sqrt(n)*qt(0.975, df=n-1)*c(-1,1)
```

```
[1] 21.52561 194.64105
```

```
set.seed(1)

meanstar = bootstrap(x=aircondit$hours, theta = mean, nboot=1000)$thetastar
seB = sd(meanstar)

# CI with bootstrap estimate of se
m + seB*qnorm(0.975)*c(-1,1)
```

```
[1] 33.69673 182.46993
```

```
m + seB*qt(0.975, df=n-1)*c(-1,1)
```



```
[1] 24.54925 191.61742
```

```
# Percetile bootstrap interval  
quantile(meanstar, c(0.025,0.975))
```

```
      2.5%      97.5%  
48.48958 190.28958
```

```
# Bootstrap BCa  
set.seed(1)  
  
bcanon(x=aircondit$hours, nboot=1000, theta=mean,  
alpha=c(0.025,0.975))$confpoints
```

```
      alpha bca point  
[1,] 0.025      57.25  
[2,] 0.975     238.50
```


```
library(bootstrap)  
  
(dif_real = mean(mouse.t) - mean(mouse.c))
```

```
[1] 30.63492
```

```
nt = length(mouse.t)  
nc = length(mouse.c)  
nperm = 1000  
set.seed(123)  
  
dif_m = vector(length=nperm)  
  
for(i in 1:nperm){  
  samp = sample(c(mouse.t, mouse.c), replace=F)  
  dif_m[i] = mean(samp[1:nt]) - mean(samp[(nt+1):(nt+nc)])  
}  
  
(ASL_perm = sum(dif_m > dif_real)/nperm)
```

```
[1] 0.144
```

```
hist(dif_m)  
abline(v=dif_real, col="violet")
```



```
t.test(mouse.t,mouse.c, alternative="greater")$p.value
```

```
[1] 0.1576861
```

One-sample randomization test

Let us test whether the median of the air-conditioning dataset is $m_0 = 6$.

There are two observations below 6 out of a total of $n = 12$ observations.

Under H_0 , the distribution of $X = \text{number of observations below 6}$ is $\text{Binomial}(n = 12, p = 0.5)$ and $P(X \leq 2) = 0.019$.

The p -value for the above test (probability of taking a sample at random with less than 2, or more than 10, observations below the median) will be thus 0.038.

```
data(aircondit, package="boot")

n = length(aircondit$hours)
(real.6 = sum(aircondit$hours < 6))
```

```
[1] 2
```

```
counter = 0
set.seed(1)

for(i in 1:10000){
  counter = counter + (sum(sample(c(0,1), n, replace=T)) <= real.6)
}

2*counter/10000
```

```
[1] 0.0368
```

The two-sample bootstrap test statistic

```
library(bootstrap)

(dif.real = mean(mouse.t) - mean(mouse.c))
```

```
[1] 30.63492
```

```
nt = length(mouse.t)
nc = length(mouse.c)
nboot = 1000
set.seed(123)

dif.m = vector(length=nboot)

for(i in 1:nboot){
  samp = sample(c(mouse.t, mouse.c), replace=T)
  dif.m[i] = mean(samp[1:nt]) - mean(samp[(nt+1):(nt+nc)])
}

(ASLboot = sum(dif.m > dif.real)/nboot)
```

```
[1] 0.127
```

```
nboot = 1000
set.seed(123)

t.hat = dif.real/(sd(c(mouse.t,mouse.c))*sqrt(1/nc+1/nt))

vt = vector(length=nboot)

for(i in 1:nboot){
  samp = sample(c(mouse.t,mouse.c),replace=T)
  dif.b = mean(samp[1:nt])-mean(samp[(nt+1):(nt+nc)])
  vt[i] = dif.b/(sd(samp)*sqrt(1/nc+1/nt))
}

(ASLt = sum(vt > t.hat)/nboot)
```

```
[1] 0.139
```

One-sample tests

```
data(aircondit)

library(MKinfer)

# Test with Ho: mu = 100
boot.t.test(aircondit[["hours"]], mu=100)
```

Bootstrap One Sample t-test

```
data:  aircondit[["hours"]]
number of bootstrap samples: 9999
bootstrap p-value = 0.7631
bootstrap mean of x (SE) = 107.7556 (34.9144)
95 percent bootstrap percentile confidence interval:
 46.58333 189.00000

Results without bootstrap:
t = 0.20554, df = 11, p-value = 0.8409
alternative hypothesis: true mean is not equal to 100
95 percent confidence interval:
 21.52561 194.64105
sample estimates:
mean of x
108.0833
```

```

unaMuestraBootpvalor = function(x, mu0, R = 1000){
  # x: observed data
  # mu0: Mean under null hypothesis

  # statistic test
  tstat = function(d, i, mu0) {
    sqrt(length(i)) * abs(mean(d[i]) - mu0) / sd(d[i])
  }

  # test statistic for observed data x
  t0 = tstat(x, 1:length(x), mu0)

  # R resampled test statistics where
  # mu0 = mean(x) passed to tstat function

  bt = boot::boot(x, tstat, R = R, mu0 = mean(x))$t[,1]

  # The p-value is obtained
  c(pvalue = mean(bt > t0))
}

```

The function is applied to a set of data in which the null hypothesis is true:

```

set.seed(666)

# H0 is correct
x = rexp(100, rate = 1/17)
unaMuestraBootpvalor(x, 17)

```

```

pvalue
0.11

```

Two-sample tests

```

DosMuestrasBootpvalor = function(x, y, alternativa = c("bilateral", "less_than", "greater_than"),
R = 2000){

  # x: observed data (first sample)
  # y: observed data (second sample)
  # alternativa - specifies the alternative hypothesis

  alternativa = match.arg(alternativa)
  n1 = length(x)
  n2 = length(y)

```

```

# test statistics
tstat = function(d, i){
  boot.xy = d[i]
  x = boot.xy[1:n1]
  y = boot.xy[-(1:n1)]
  s = sqrt( ((n1-1) * var(x) + (n2 - 1) * var(y)) / (n1 + n2 - 2) )
  mu.x = mean(x)
  mu.y = mean(y)
  (mu.x - mu.y) / s / sqrt(1 / n1 + 1 / n2)
}

xy = c(x,y)

# Test statistics for the observed data
t0 = tstat(xy, 1:(n1 + n2))

# Resampling of the test statistic
bt = boot::boot(xy, tstat, R = R)$t[,1]

# p-value
if(alternativa == "greater_than") return(c(pvalue = mean(bt > t0)))
if(alternativa == "less_than") return(c(pvalue = mean(bt < t0)))
c(pvalue = mean(abs(bt) > abs(t0)))
}

```

The function is applied to a set of data in which the null hypothesis is true:

```

# H0 is correct, mu_x is equal to mu_y
set.seed(666)
datos1 = rnorm(10, mean = 3, sd = 2)
datos2 = rnorm(20, mean = 3, sd = 2)

DosMuestrasBootpvalor(datos1, datos2, alternativa = "greater_than")

```

```

pvalue
0.313

```

```

# library(MKinfer)

boot.t.test(datos1, datos2)

```

Bootstrap Welch Two Sample t-test

```

data:  datos1 and datos2
number of bootstrap samples: 9999
bootstrap p-value = 0.6675
bootstrap difference of means (SE) = 0.4797342 (1.008988)
95 percent bootstrap percentile confidence interval:

```

```
-1.537767  2.465471
```

Results without bootstrap:

t = 0.44627, df = 14.379, p-value = 0.662

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

```
-1.804340  2.755463
```

sample estimates:

mean of x mean of y

```
2.807849  2.332287
```

It can be compared with a permutation test:

```
# library(MKinfer)
```

```
perm.t.test(datos1, datos2)
```

Permutation Welch Two Sample t-test

data: datos1 and datos2

number of permutations: 9999

(Monte-Carlo) permutation p-value = 0.6545

permutation difference of means (SE) = 0.4712101 (0.9550028)

95 percent (Monte-Carlo) permutation percentile confidence interval:

```
-1.349719  2.308100
```

Results without permutation:

t = 0.44627, df = 14.379, p-value = 0.662

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

```
-1.804340  2.755463
```

sample estimates:

mean of x mean of y

```
2.807849  2.332287
```